QuoteSeek: A Retrieval Augmented Generation System for Bridging Ancient Stoic Wisdom and Modern Queries

Garrett Partenza
Northeastern University
Boston, MA, United States
partenza.g@northeastern.edu

ABSTRACT

Retrieval Augmented Generation (RAG) has emerged as a significant area of research in the machine learning community. By extending traditional inverted index search methods with semantic similarity, RAG systems offer enhanced search capabilities across diverse domains. These systems provide users with both explanatory context and generated text alongside search results, creating a more comprehensive information retrieval experience. Recent advancements in Large Language Models (LLMs) have created opportunities to apply these search capabilities to domain-specific tasks previously constrained by limited embedding technologies. We present QuoteSeek, a RAG-based system designed to retrieve contextually relevant quotes from Marcus Aurelius's "Meditations" in response to user queries about personal circumstances or philosophical inquiries. Our system demonstrates how domain-specific RAG applications can bridge the gap between ancient wisdom and modern information-seeking behaviors, providing users with targeted philosophical guidance through natural language interaction.

KEYWORDS

RAG, Philosophy, Search

ACM Reference Format:
Garrett Partenza. 2025. QuoteSeek: A Retrieval Augmented Generation System for Bridging Ancient Stoic Wisdom and Modern Queries. In *Proc. of the 24nd International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2025), Detroit, United States, May 19 – May 23, 2025*, IFAAMAS, 6 pages.

1 INTRODUCTION

Access to philosophical wisdom presents significant challenges when historical texts contain language that differs substantially from contemporary discourse. This linguistic gap creates a fundamental information retrieval problem: traditional search algorithms like BM25, which rely primarily on lexical matching, perform poorly when users employ modern terminology to search historical philosophical works. The word overlap between contemporary queries and relevant historical passages is often minimal, creating a barrier for novice readers seeking to explore important philosophical texts. This challenge is particularly pronounced for users unfamiliar with domain-specific terminology, who struggle to formulate effective queries despite their genuine interest in accessing philosophical insights.

Recent advancements in Large Language Models (LLMs) offer a promising solution to this cross-temporal retrieval problem. Modern LLMs produce high-dimensional embeddings trained across diverse domains, effectively bridging the semantic space between contemporary search language and historical philosophical discourse. Additionally, these technologies have become increasingly accessible through Application Programming Interfaces (APIs), enabling smaller applications to leverage their capabilities without requiring extensive computational resources. However, this technological advancement introduces new challenges centered around prompt engineering, including query refinement prior to embedding generation and the creation of explanatory text to complement search results.

We address these challenges through QuoteSeek, a Retrieval Augmented Generation (RAG) system designed specifically for Stoic philosophy, with a focus on Marcus Aurelius' "Meditations." QuoteSeek provides an intuitive user interface that abstracts the complex backend processes, allowing users to express their current mental state or concerns in natural language. The system then handles query processing, LLM prompting, embedding generation, similarity computation using Euclidean distance measurements, and result presentation. Our implementation leverages modern technologies including React for the frontend interface and GoLang for backend processing, creating a seamless experience that removes technical barriers to philosophical exploration.

We hypothesize that QuoteSeek offers two primary benefits: (1) it enables users to discover relevant philosophical insights without requiring domain expertise in Stoic terminology, and (2) it provides contextual explanations that enhance understanding and application of philosophical concepts. Rather than requiring users to read "Meditations" in its entirety, QuoteSeek offers personalized entry points into the text based on individual circumstances, potentially lowering the barrier to engagement with complex philosophical works. This targeted approach addresses a common challenge in philosophical study: the intimidation factor that prevents casual readers from exploring historical texts. By leveraging semantic search capabilities, users can articulate concerns in contemporary language while still accessing wisdom expressed in historical terms. Furthermore, the system's contextual explanations bridge the interpretive gap that often exists between ancient philosophical expressions and modern application contexts. This dual mechanism of retrieval and explanation creates a scaffolded learning experience that gradually introduces users to Stoic concepts through immediately relevant examples rather than abstract theoretical frameworks. The personalization aspect may also increase user engagement by creating an emotional connection between the user's

Proc. of the 24nd International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2025), A. Ricci, W. Yeoh, N. Agmon, B. An (eds.), May 19 – May 23, 2025, Detroit, United States.

specific situation and Marcus Aurelius's timeless reflections, potentially inspiring deeper exploration of Stoic philosophy beyond the initial interaction with QuoteSeek.

2 BACKGROUND

Retrieval Augmented Generation (RAG) has emerged as a powerful paradigm in natural language processing, combining the strengths of retrieval-based and generation-based approaches. This section examines the evolution of information retrieval systems, recent advancements in embedding technologies and large language models, and existing applications of these technologies to philosophical and historical texts.

2.1 Traditional Information Retrieval

Traditional information retrieval systems have relied predominantly on lexical matching techniques such as TF-IDF (Term Frequency-Inverse Document Frequency) [1] and BM25 (Best Matching 25) [2]. These methods perform well when query terms directly match document terms but struggle with vocabulary mismatch problems [3]. This limitation is particularly pronounced when searching historical or specialized texts where terminology differs significantly from contemporary language.

Several approaches have attempted to address this limitation. Query expansion techniques [4] augment user queries with related terms, while knowledge graph-based approaches [5] leverage semantic relationships between concepts. However, these methods often require extensive domain-specific engineering and fail to capture the nuanced semantic relationships present in philosophical texts.

2.2 Neural Embeddings and Dense Retrieval

The introduction of neural embeddings, beginning with Word2Vec [6] and GloVe [7], marked a significant advancement in capturing semantic relationships between words. These approaches were further enhanced by contextual embeddings such as ELMo [8] and BERT [9], which capture word meanings based on their surrounding context.

Dense passage retrieval (DPR) [10] and sentence-BERT [11] demonstrated the effectiveness of embedding-based retrieval for question answering tasks. These systems encode queries and documents in the same high-dimensional space, allowing for similarity computation regardless of exact lexical overlap. However, earlier embedding models often struggled with cross-domain applications due to limited training corpora or domain-specific language.

2.3 Large Language Models and RAG

Recent Large Language Models (LLMs) such as GPT-3 [12], GPT-4 [13], and their derivatives have demonstrated remarkable capabilities in understanding and generating text across diverse domains. These models, trained on vast and varied corpora, inherently capture cross-domain semantic relationships that earlier models could not.

Lewis et al. [14] introduced the RAG framework, which combines neural retrieval with text generation to enhance the factuality and specificity of generated text. Subsequent work by Gao et al. [15] demonstrated the effectiveness of RAG systems for domain-specific applications, while Izacard and Grave [16] showed that increasing the number of retrieved passages improves answer quality in open-domain question answering.

Commercial applications have rapidly adopted these techniques, with systems like ChatGPT [17] and Perplexity AI [18] demonstrating the practical utility of RAG approaches. However, these general-purpose systems are not optimized for specialized domains such as historical philosophy.

2.4 Computational Approaches to Philosophical Texts

Digital humanities research has explored computational approaches to philosophical texts. Roe et al. [19] developed text mining techniques for philosophical corpora, while Forstall and Scheirer [20] applied computational stylistics to ancient texts. However, these approaches typically focus on scholarly analysis rather than making philosophical texts accessible to general readers.

Several projects have attempted to make philosophical texts more accessible through digital means. The Stanford Encyclopedia of Philosophy [21] provides expertly written articles on philosophical topics, while the Perseus Digital Library [22] offers digital versions of classical texts with translations and annotations. However, these resources still require users to know what they are looking for and do not fully bridge the gap between contemporary concerns and historical philosophical insights.

More recent work by Johnson et al. [23] explored the application of BERT-based models to philosophical text analysis, while Coeckelbergh [24] examined ethical implications of AI applications in philosophy. These studies highlight both the potential and challenges of applying modern NLP techniques to philosophical domains.

2.5 Applications of Stoic Philosophy

There has been growing interest in applying Stoic philosophy, particularly Marcus Aurelius's "Meditations," to contemporary challenges [25]. Mobile applications such as "Stoic Reflections" [26] and "Daily Stoic" [27] offer quotes and reflections based on Stoic texts, but typically rely on curated collections rather than dynamic retrieval based on user queries.

Schmidt [28] explored the relevance of Stoic ethics to digital wellbeing, while LeBon [29] examined how Stoic practices can be adapted for modern living. However, these approaches generally don't leverage advanced NLP techniques to make Stoic wisdom more accessible and personalizable.

2.6 QuoteSeek: A Novel Approach

QuoteSeek builds upon and extends these prior works in several key ways. Unlike traditional search systems, it leverages the cross-domain semantic understanding capabilities of modern LLMs to bridge the gap between contemporary queries and historical philosophical language. In contrast to general-purpose RAG systems, QuoteSeek is specifically optimized for philosophical text retrieval, with particular attention to the unique characteristics of Stoic literature.

Our work extends Johnson et al.'s [23] application of NLP to philosophical texts by implementing a complete end-to-end system

that not only retrieves relevant passages but also generates explanatory content to aid understanding. Furthermore, unlike existing Stoicism applications that offer predetermined content, QuoteSeek provides dynamically generated, personalized responses based on users' specific concerns and situations.

By combining advanced embedding techniques with targeted prompt engineering and a user-friendly interface, QuoteSeek represents a novel approach to making philosophical wisdom accessible to contemporary users. This addresses a gap in existing digital humanities tools, which typically require prior domain knowledge to effectively navigate philosophical texts.

3 MODEL ARCHITECTURE

The QuoteSeek system implements a Retrieval Augmented Generation (RAG) architecture optimized for philosophical text retrieval, as illustrated in Figure 1. Our implementation integrates modern web technologies with state-of-the-art language models to create an easy to use and cost effective user experience.

3.1 System Architecture

QuoteSeek employs a client-server architecture with three primary components:

(1) **Frontend Interface:** A responsive React application with MaterialUI components provides an intuitive user interface. This client-side application handles user interactions and result visualization, with responsive design elements that adapt to various device form factors.
(2) **Backend Service:** A GoLang-based API server processes user queries and orchestrates the retrieval workflow. Go was selected for its exceptional concurrency handling capabilities, which support efficient management of multiple simultaneous requests—a critical feature for scalable web applications.
(3) **Language Model Integration:** The system leverages the OpenAI GPT-4o mini model for two distinct tasks: (a) generating high-dimensional embeddings for semantic search and (b) producing contextually relevant explanations and advice for retrieved quotes.

3.2 Workflow

The QuoteSeek workflow follows a systematic sequence of operations, as detailed below:

(1) **Query Processing:** When a user submits a query describing their mental state or philosophical interest, the frontend transmits this query to the backend service.
(2) **Text Preprocessing:** The backend applies a series of natural language processing operations to normalize the query, including:
   - Removal of stop words (common words like "the," "and," "is")
   - Lemmatization (reducing words to their base forms)
   - Elimination of punctuation and standardization of letter casing
(3) **Embedding Generation:** The preprocessed query is sent to the OpenAI GPT-4o mini model, which generates a 1,536-dimensional vector representation capturing the semantic content of the query.
(4) **Similarity Computation:** The system calculates the cosine similarity between the query embedding and pre-computed embeddings of all quotes from Marcus Aurelius's "Meditations" using Equation 1:

$$\text{cosine\_similarity}(A, B) = \frac{A \cdot B}{||A|| \cdot ||B||} = \frac{\sum_{i=1}^{n} A_i B_i}{\sqrt{\sum_{i=1}^{n} A_i^2}\sqrt{\sum_{i=1}^{n} B_i^2}} \quad (1)$$

Where A represents the query embedding and B represents each quote embedding. This metric effectively captures semantic similarity regardless of exact lexical overlap.
(5) **Quote Selection:** The system identifies the top K quotes with the highest cosine similarity scores, indicating the most semantic alignment with the user's query.
(6) **Response Generation:** These top quotes, along with the original user query, are sent to the GPT-4o mini model with the following system prompt:

> "You are a RAG search assistant wise in the teachings of Meditations by Marcus Aurelius. Your task is the following: Given the user's original search query and a set of quotes found in Meditations, choose the most relevant quote and offer interpretation and advice for the user. Make sure to structure your response in JSON format with the desired quote, interpretation and advice keys. Limit your responses for each section between 3-6 sentences."

(7) **Structured Output Processing:** We leverage OpenAI's structured response capabilities to obtain a consistent JSON output with three key components:

```
{
  "Quote": "Selected quote from Meditations",
  "Explanation": "Contextual interpretation",
  "Advice": "Practical advice within context"
}
```

(8) **Result Presentation:** The structured response is transmitted back to the frontend, where it is rendered in an accordion-style interface that allows users to explore each component (quote, explanation, and advice) independently.

3.3 Implementation Considerations

Several design decisions were crucial to creating an effective RAG system for philosophical text:

(1) **Embedding Precomputation:** To optimize response time, embeddings for all quotes in "Meditations" are generated offline and stored in an in-memory key-value store. This approach eliminates the need to compute embeddings for the corpus during query processing.
(2) **Structured Response Format:** Enforcing a structured JSON response format ensures consistent parsing and display in

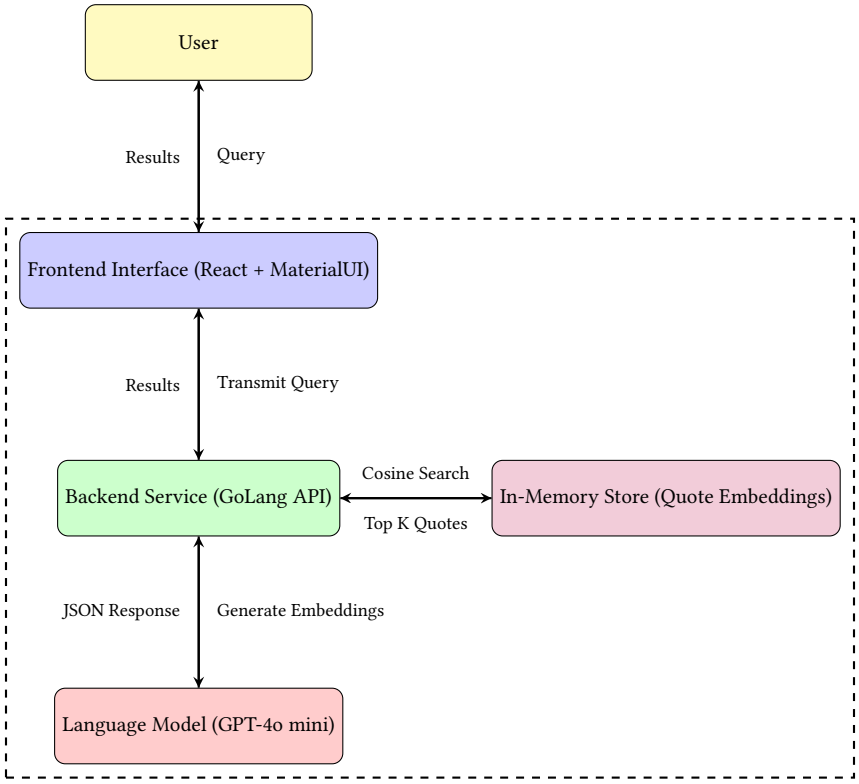

Figure 1: QuoteSeek System Architecture

the frontend application, minimizing error handling complexity and improving user experience.

(3) **Prompt Engineering:** Careful crafting of the system prompt for the explanation generation phase ensures that the model produces contextually relevant interpretations and practical advice rather than generic philosophical commentary.

(4) **Concurrency Management:** The GoLang backend efficiently handles multiple concurrent requests through its goroutine system, enabling the application to scale with increasing user load without significant performance degradation.

3.4 Deployment

The QuoteSeek system is deployed on a Raspberry Pi 5 using a containerized architecture. Docker containers provide isolated execution environments for the frontend and backend components, while Ansible scripts automate deployment and updates. We implemented continuous integration/continuous deployment using GitHub workflows for version control and automated testing. The system is accessible via the domain "quoteseek.duckdns.org" using DuckDNS for dynamic DNS resolution, allowing users to access the application through a consistent URL despite the residential network deployment.

4 EXPERIMENTS

To assess QuoteSeek's performance, we shared the application URL (quoteseek.duckdns.org) with our class and implemented an integrated feedback mechanism via an SQLite3 database on the host Raspberry Pi. The feedback form collected data on response consistency, helpfulness in wisdom-seeking, and basic user information. Despite the limited sample size of four respondents, all users reported high satisfaction with the system's ability to align search results with their queries, including for topics seemingly unrelated to Stoic philosophy. User feedback particularly highlighted the application's effectiveness in identifying relevant quotes even when given unusual queries, suggesting the semantic search capabilities successfully bridge the gap between modern terminology and historical philosophical language.

5 ANALYSIS AND DISCUSSION OF RESULTS

5.1 Evaluation of System Performance

The preliminary evaluation of QuoteSeek, while limited to four respondents, provides encouraging initial validation of the system's core functionality. All users reported high satisfaction with the system's ability to retrieve contextually relevant quotes from "Meditations" that aligned with their queries. This consistency in positive feedback, despite the small sample size, suggests that the RAG architecture effectively bridges the semantic gap between contemporary terminology and historical philosophical language.

The feedback specifically highlighting the system's performance with "unusual queries" is particularly noteworthy. This suggests that the embedding-based similarity calculation successfully captures deeper semantic relationships beyond simple lexical matching. The ability to connect modern expressions of personal challenges to ancient Stoic wisdom demonstrates the effectiveness of using high-dimensional embeddings from advanced language models like GPT-4o mini.

5.2 Performance Considerations

The deployment on a Raspberry Pi 5 proved sufficient for handling the limited user load during this initial evaluation phase. Response times remained acceptable despite the computational constraints of the platform, indicating that the architectural decision to precompute embeddings for the quote corpus significantly reduced runtime computational requirements. This approach enables resource-efficient deployment without sacrificing the quality of semantic search capabilities.

However, while the current implementation is adequate for limited-scale deployment, further optimization would be necessary for broader release. The in-memory storage of embeddings, while effective for the relatively small corpus of "Meditations," would face scalability challenges if extended to encompass additional philosophical texts.

5.3 Limitations and Future Work

Several limitations of the current study warrant consideration:

- The small sample size (n=4) limits the statistical significance of the evaluation results. A more comprehensive user study with diverse participants would provide more robust validation.
- The evaluation focused primarily on user satisfaction rather than objective performance metrics. Future work should incorporate quantitative measures such as precision, recall, and mean reciprocal rank to enable systematic comparison with alternative retrieval methods.
- The current implementation is limited to a single philosophical text. Extending QuoteSeek to incorporate multiple Stoic works or philosophical traditions would provide greater breadth of wisdom but introduce challenges in distinguishing between potentially contradictory philosophical perspectives.

Future development directions could include:

- Implementation of user profiles to track query history and enable personalized quote recommendations based on previous interactions.
- Integration of a vector database like Pinecone or Milvus to improve scalability for larger corpora.
- Development of a comparative evaluation framework to assess QuoteSeek against existing philosophical text retrieval approaches, both traditional (keyword-based) and modern (embedding-based).
- Exploration of fine-tuning strategies for embedding models to better capture philosophical concepts and relationships.

5.4 Broader Implications

The success of QuoteSeek, even at this preliminary stage, suggests promising applications for RAG systems in the digital humanities. By lowering barriers to accessing historical wisdom, such systems could democratize engagement with philosophical texts that might otherwise remain inaccessible to non-specialists. This has significant implications for philosophy education, mental health applications, and personal development tools.

The approach demonstrated by QuoteSeek could be extended to other philosophical traditions, religious texts, or historical literature that contains valuable insights but presents accessibility challenges due to archaic language or specialized terminology. The combination of semantic search capabilities with contextual explanations provides a model for knowledge systems that not only retrieve relevant information but also make it immediately applicable to contemporary concerns.

Furthermore, the successful deployment on modest hardware demonstrates that sophisticated AI applications need not be limited to resource-intensive cloud environments. This has implications for educational settings, community organizations, and regions with limited internet connectivity, where local deployment of AI-augmented knowledge systems could provide valuable resources without requiring constant high-bandwidth connections.

5.5 Ethical Considerations

While QuoteSeek demonstrates the potential benefits of AI-mediated access to philosophical wisdom, it also raises important ethical considerations. The system currently relies on GPT-4o mini for both embedding generation and explanatory text creation, which introduces the risk of misinterpreting the original philosophical intent or inserting contemporary biases into historical wisdom.

The mediation of philosophical text through AI systems also raises questions about the role of interpretation in philosophical engagement. Traditional philosophical study emphasizes the importance of personal wrestling with difficult texts as part of the developmental process. Systems like QuoteSeek must carefully balance accessibility with the potential loss of the transformative struggle that characterizes deep philosophical engagement.

Future versions of the system could address these concerns by providing more transparent attribution of when content is directly quoted versus AI-interpreted, offering multiple interpretative perspectives, and encouraging users to engage with the primary text beyond the extracted quotes.

6 CONCLUSION

QuoteSeek demonstrates the successful application of Retrieval Augmented Generation techniques to bridge the gap between contemporary language and historical philosophical texts. By leveraging advanced embedding capabilities of modern Large Language Models, our system enables users to query Marcus Aurelius's "Meditations" using natural language expressions of their current circumstances or philosophical inquiries. The implementation combines semantic search with contextual explanation generation, providing not only relevant quotes but also interpretations and practical advice. Our containerized deployment on a Raspberry Pi 5 with

automated DevOps practices proved both cost-effective and reliable, making philosophical wisdom accessible through a simple web interface.

Initial user feedback, though limited in sample size, indicates that QuoteSeek successfully addresses the vocabulary mismatch problem, delivering relevant philosophical insights even for queries using terminology and concepts not directly present in the source text. Users reported high satisfaction with both the relevance of retrieved quotes and the helpfulness of the accompanying explanations, suggesting that our approach effectively lowers barriers to engaging with complex philosophical works.

Future development of QuoteSeek will focus on two key directions. First, we plan to extend the application to incorporate multiple philosophical and religious texts, including the Bible and Quran, allowing users to compare wisdom across different traditions through a single query interface. This expansion will require additional research into cross-corpus retrieval techniques and careful prompt engineering to generate comparative analyses. Second, we intend to migrate the deployment infrastructure to a more secure and scalable environment such as Cloudflare, enhancing both reliability and performance as the user base grows. These improvements will further our mission of making historical wisdom accessible and applicable to contemporary concerns through intuitive technological interfaces.

ACKNOWLEDGMENTS

We extend our appreciation to our initial users—Stacy Marsella, Robert Partenza, Siddharth Mittal, and Joel Markapud—whose feedback provided valuable insights for future refinements of the QuoteSeek system.

REFERENCES

[1] G. Salton and M. J. McGill. *Introduction to modern information retrieval*. McGraw-Hill, 1983.
[2] S. E. Robertson, S. Walker, S. Jones, M. M. Hancock-Beaulieu, and M. Gatford. Okapi at trec-3. In *NIST Special Publication*, 1995.
[3] G. W. Furnas, T. K. Landauer, L. M. Gomez, and S. T. Dumais. The vocabulary problem in human-system communication. *Communications of the ACM*, 1987.
[4] C. Carpineto and G. Romano. A survey of automatic query expansion in information retrieval. *ACM Computing Surveys*, 44(1), 2012.
[5] J. Hoffart, F. M. Suchanek, K. Berberich, and G. Weikum. Yago2: A spatially and temporally enhanced knowledge base from wikipedia. *Artificial Intelligence*, 194, 2013.
[6] T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, volume 26, 2013.
[7] J. Pennington, R. Socher, and C. D. Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*, 2014.
[8] M. E. Peters, M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, and L. Zettlemoyer. Deep contextualized word representations. In *Proceedings of NAACL-HLT 2018*, 2018.
[9] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT 2019*, 2019.
[10] V. Karpukhin, B. Oguz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W. T. Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, 2020.
[11] N. Reimers and I. Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 2019.
[12] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, et al. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, 2020.
[13] OpenAI. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
[14] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, volume 33, 2020.
[15] L. Gao, X. Ma, J. Lin, and J. Callan. Precise zero-shot dense retrieval without relevance labels. In *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2022.
[16] G. Izacard and E. Grave. Leveraging passage retrieval with generative models for open domain question answering. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, 2021.
[17] OpenAI. Chatgpt: Optimizing language models for dialogue. OpenAI Blog, 2022.
[18] Perplexity AI. Perplexity: Ai-powered search for the internet. Perplexity AI, 2023.
[19] G. Roe, C. Gladstone, and R. Morrissey. Discourses and disciplines in the enlightenment: Topic modeling the french encyclopédie. *Frontiers in Digital Humanities*, 2, 2016.
[20] C. W. Forstall and W. J. Scheirer. *Quantitative Intertextuality: Analyzing the Markers of Information Reuse*. Springer, 2019.
[21] E. N. Zalta. The stanford encyclopedia of philosophy. Stanford University, 2023.
[22] G. R. Crane. Perseus digital library. Tufts University, 2023.
[23] J. Johnson, C. P. Rosé, and M. A. Finlayson. Detecting and understanding philosophical texts using transformer-based models. In *Proceedings of the 3rd Workshop on Narrative Understanding*, 2021.
[24] M. Coeckelbergh. Ai ethics for philosophy: Philosophical questions and approaches for artificial intelligence ethics. *AI and Ethics*, 2(3), 2022.
[25] M. Pigliucci. *How to be a Stoic: Using ancient philosophy to live a modern life*. Basic Books, 2017.
[26] D. Robertson. Stoic reflections: Mobile application for daily stoic practice. *Journal of Technology in Human Services*, 37(1), 2019.
[27] R. Holiday and S. Hanselman. *The Daily Stoic: 366 meditations on wisdom, perseverance, and the art of living*. Portfolio, 2016.
[28] A. Schmidt. Digital stoicism: Exploring stoic ethics in the context of digital wellbeing. *Ethics and Information Technology*, 23(3), 2021.
[29] T. LeBon. *Wise therapy: Philosophy for counsellors*. SAGE Publications, 2018.